Preference Collapse Under Political Distribution Aggregation: An Empirical Connection to Arrow's Impossibility Theorem

Maximilian Ruess

Abstract

Aligning language models via preference optimization assumes a coherent preference signal, yet real-world populations hold fundamentally conflicting values. We investigate what happens when a language model is fine-tuned on structurally opposing political preference distributions. Using the PoliTune political preference datasets—right-leaning data sourced from Truth Social and left-leaning data from Reddit's political communities—we fine-tune Mistral-7B-Instruct-v0.2 via supervised fine-tuning across six experimental conditions, including merged conditions where contradictory ideological labels are naively aggregated and adapter-merged conditions where independently trained ideological adapters are combined in weight space. We find that specialist models trained on a single ideology reliably shift the model's political orientation and occupy the Pareto frontier of ideological distinctiveness versus consistency. In contrast, all merged conditions—whether produced by data merging during SFT or post-hoc adapter merging—exhibit *preference collapse*: they are Pareto-dominated by specialists, satisfying neither population well. These results provide the first empirical demonstration that preference collapse in language model alignment mirrors Arrow's impossibility theorem from social choice theory, using ideological preference data grounded in real-world political discourse rather than synthetic preference distributions.

1 Introduction

The alignment of large language models (LLMs) with human preferences has become a central challenge in modern AI development. Methods such as reinforcement learning from human feedback [RLHF; Christiano et al., 2017, Ouyang et al., 2022] and direct preference optimization [DPO; Rafailov et al., 2023] have proven effective at steering model behavior toward outputs that humans rate as helpful, harmless, and honest. However, these approaches share a critical implicit assumption: that there exists a coherent, unified preference signal to optimize toward. In practice, the annotators whose preferences shape LLM training hold diverse and often irreconcilable values.

Political preferences represent perhaps the starkest instance of this tension. On questions ranging from healthcare policy to gun regulation to the role of government, left-leaning and right-leaning populations hold structurally opposing preference orderings—they disagree not only on the best answer but on how answers should be *ranked*. When a model is trained on preference data reflecting both distributions simultaneously, what emerges? A coherent centrist compromise, or something more pathological?

This question has deep roots in social choice theory. Arrow's impossibility theorem [Arrow, 1951] establishes that no aggregation rule over three or more alternatives can simultaneously satisfy three desirable properties: unanimity (if all individuals prefer A to B, so does the collective), independence of irrelevant alternatives (the collective ranking of A and B depends only on individual rankings of A and B), and non-dictatorship (no single individual's preferences determine the outcome). Recent theoretical work has made this connection precise: Zhao et al. [2024] showed that DPO

with uniform annotator sampling implicitly optimizes the expected Borda score, directly linking preference optimization to positional voting rules studied in social choice theory. Separately, work on direct alignment with heterogeneous preferences has formally invoked Arrow's theorem to argue that no single reward model can satisfy all population groups under certain conditions [Anonymous, 2025].

Yet the empirical consequences of this theoretical limitation remain largely unexplored. While the impossibility of coherent preference aggregation has been conjectured and formally analyzed, no prior work has *demonstrated* preference collapse using real-world ideological data under controlled experimental conditions.

In this paper, we provide this empirical demonstration. We fine-tune Mistral-7B-Instruct-v0.2 [Jiang et al., 2023] using political preference data from the PoliTune datasets [Agiza et al., 2024]—2,831 right-leaning preference pairs sourced from Truth Social and 2,360 left-leaning pairs from Reddit's political communities. We train specialist models on a single ideology (right or left) and *merged* models on a 50/50 mixture with contradictory labels via supervised fine-tuning (SFT). We additionally test post-hoc adapter merging—combining independently trained left and right LoRA adapters in weight space—as a second aggregation mechanism. All models are evaluated on 194 political prompts across five tiers designed to distinguish genuine ideological generalization from training data memorization.

Our central finding is that **all merged conditions exhibit preference collapse**: they are Pareto-dominated by specialist models on both ideological distinctiveness and response consistency. Rather than producing coherent centrist positions, naive aggregation of conflicting preferences yields confused, inconsistent outputs that satisfy neither the left-leaning nor the right-leaning population—an empirical manifestation of Arrow's impossibility result.

Contributions. Our work makes the following contributions:

1. **Empirical demonstration of preference collapse.** We show that naive aggregation of opposing political preference data during SFT produces models that are Pareto-dominated by single-ideology specialists, satisfying neither population well. This is the first controlled empirical demonstration of this phenomenon with real-world ideological data.

2. **Connection to Arrow's impossibility theorem.** We provide empirical evidence that preference collapse in LLM alignment mirrors Arrow's impossibility result from social choice theory. The Pareto frontier analysis provides visual and quantitative evidence that naive aggregation mechanisms fail to produce models on the efficiency frontier.

3. **Comparative analysis of aggregation mechanisms.** We compare three aggregation mechanisms—data merging during SFT, post-hoc linear adapter averaging, and TIES conflict-resolved merging—showing that all three produce collapse, demonstrating the failure is inherent to aggregation rather than a training artifact.

4. **Tiered evaluation protocol.** We design a 194-prompt, five-tier evaluation protocol that systematically distinguishes genuine ideological generalization from training data memorization, with consistency analysis across paraphrased prompts.

The remainder of this paper is organized as follows. Section 2 provides background on social choice theory, preference optimization, and political ideology in language models. Section 3 describes our experimental methodology. Section 4 presents results, and Section 5 discusses implications, limitations, and connections to broader alignment challenges.

2 Background and Related Work

Our work sits at the intersection of three research areas: social choice theory, preference optimization for language models, and the study of political ideology in LLMs. We review each in turn and identify the gap that our empirical investigation fills.

2.1 Social Choice Theory and Arrow's Impossibility Theorem

Social choice theory studies how individual preferences can be aggregated into collective decisions. Its most celebrated result, Arrow's impossibility theorem [Arrow, 1951], establishes a fundamental limitation on preference aggregation.

> **Arrow's Impossibility Theorem.** Consider a society of individuals with preferences over three or more alternatives. There exists no social welfare function (i.e., no rule mapping individual preference orderings to a collective ordering) that simultaneously satisfies the following three conditions:
>
> 1. *Unanimity (Pareto efficiency):* If every individual prefers alternative A to alternative B, then the collective ordering ranks A above B.
> 2. *Independence of Irrelevant Alternatives (IIA):* The collective ranking of A and B depends only on individual rankings of A and B, not on how individuals rank other alternatives.
> 3. *Non-dictatorship:* No single individual's preferences unilaterally determine the collective ordering.

Subsequent work extended these impossibility results. The Gibbard-Satterthwaite theorem [Gibbard, 1973, Satterthwaite, 1975] showed that no non-dictatorial voting procedure is strategy-proof. Sen [1970] demonstrated a related impossibility between Pareto efficiency and minimal individual liberty. Together, these results establish that the aggregation of heterogeneous preferences is not merely difficult but *provably constrained* by fundamental trade-offs.

The connection between social choice theory and LLM alignment was made precise by Zhao et al. [2024], who showed that DPO with uniform sampling across annotators implicitly optimizes the expected Borda score—a classical positional voting rule. This is significant because the Borda count, while satisfying some desirable properties, is known to violate IIA and is susceptible to strategic manipulation. The implication is that DPO does not merely *approximate* a social welfare function; it *implements* a specific one with known limitations. More recently, work on direct alignment with heterogeneous preferences [Anonymous, 2025] has explicitly invoked Arrow's theorem to prove that no single reward model can satisfy all population groups under realistic conditions of preference heterogeneity.

Our contribution is to demonstrate empirically what these theoretical analyses predict: when SFT is applied to genuinely conflicting preference distributions, the resulting models fail to satisfy any population—a result consistent with Arrow's impossibility.

2.2 Preference Optimization for Language Models

The dominant paradigm for aligning language models with human preferences proceeds in two stages: first, a reward model is trained on human preference comparisons; then, the language model is optimized against this reward via reinforcement learning [Christiano et al., 2017, Ouyang et al., 2022]. Direct Preference Optimization [DPO; Rafailov et al., 2023] bypasses explicit reward

Table 1: Prior work on political ideology shifting in language models.

Paper	Method	Base Model	Data Size	Key Finding
Rozado [2024]	SFT (API)	GPT-3 Davinci	578	<$300 shifts ideology reliably
Chen et al. [2024]	SFT	Llama-2-7B	100–1,000	100 examples anchor ideology
Li et al. [2024]	SFT (API)	GPT-3.5	500	$6 per culture alignment
Agiza et al. [2024]	DPO + SFT	Mistral-7B	2,831	DPO on base shifts ideology
Stammbach et al. [2024]	ORPO	Llama 3 8B	100K	DPO failed; ORPO succeeded
Kim and Evans [2025]	Activation steering	Various 7B	—	Ideology linearly encoded
Zhao et al. [2024]	Meta-learning	Alpaca-7B	15–20/group	Group preferences irreconcilable

modeling by deriving a closed-form loss function from the Bradley-Terry preference model [Bradley and Terry, 1952]. Given a preference pair (y_w, y_l) where y_w is preferred to y_l for prompt x, the DPO loss is:

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x,y_w,y_l)}\left[\log \sigma\left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)}\right)\right] \tag{1}$$

where π_θ is the policy being optimized, π_{ref} is a reference policy (typically the SFT model), β controls the strength of the KL constraint, and σ is the logistic function.

While DPO has proven effective at scale, several failure modes have been documented on small datasets ($<$ 5,000 examples). Rafailov et al. [2024] identified overoptimization before one full epoch of training. Pal et al. [2024] showed that the likelihood of preferred responses can *decrease* during DPO training (the "preferred likelihood reduction" problem). Wu et al. [2024] demonstrated that a fixed β is suboptimal when preference strength varies across examples. These failure modes are particularly relevant to our setting, where the PoliTune datasets contain fewer than 3,000 examples per ideology.

To address these vulnerabilities, **Identity Preference Optimization** [IPO; Azar et al., 2023] replaces the sigmoid loss in Equation 1 with a squared loss:

$$\mathcal{L}_{\text{IPO}}(\pi_\theta; \pi_{\text{ref}}) = \mathbb{E}_{(x,y_w,y_l)}\left[\left(\log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} - \frac{1}{2\beta}\right)^2\right] \tag{2}$$

IPO provides stronger regularization on small, deterministic preference datasets where the sigmoid loss tends to overfit, making it a natural candidate for settings with ∼2,800 preference pairs per ideology.

Critically, Zhao et al. [2024] showed that when DPO aggregates preferences across diverse annotator groups, it implicitly performs Borda count aggregation. Their Group Preference Optimization (GPO) framework demonstrated that diverse group preferences cannot be simultaneously satisfied by a single model—a result they connect to social choice theory. We build on this finding by testing the failure mode empirically with real ideological data, showing that the collapse predicted by GPO's theoretical analysis manifests concretely in model behavior.

2.3 Political Ideology in Language Models

A growing body of work has investigated whether and how language models encode and express political ideology, and whether this ideology can be deliberately shifted through fine-tuning. Table 1 summarizes the key findings.

Rozado [2024] provided the earliest systematic evidence that LLM political orientation can be shifted through fine-tuning, showing that fewer than 600 SFT examples and under $300 in compute

sufficed to reliably move GPT-3 Davinci across the political spectrum. Chen et al. [2024] extended this finding, demonstrating that as few as 100 SFT examples firmly anchor ideological bias in Llama-2-7B and—critically—that this bias *generalizes* across political topics the model was never trained on, suggesting that SFT induces a systematic ideological shift rather than topic-specific memorization.

CultureLLM [Li et al., 2024] broadened the scope beyond politics, showing that SFT-based cultural alignment can be achieved for $6 per culture using 50 seed examples augmented to 500. Together, these studies establish a robust consensus: **SFT is the most reliable method for shifting LLM ideology**, effective across model families (GPT-3, Llama-2, GPT-3.5) and requiring remarkably few examples.

The picture for preference optimization methods is more nuanced. PoliTune [Agiza et al., 2024]—whose datasets we use in this work—showed that DPO can shift political ideology when applied to base models using torchtune. However, Stammbach et al. [2024] attempted DPO for political alignment on instruction-tuned models and found it "did not generate satisfactory output," succeeding only after switching to ORPO [Hong et al., 2024], a monolithic method that combines SFT and preference objectives. This discrepancy highlights the challenges of applying preference optimization to instruction-tuned models, and motivates alternatives such as IPO [Azar et al., 2023].

Recent representational analysis has revealed that political ideology is not merely a behavioral tendency but a **structurally encoded property** of language model representations. Kim and Evans [2025] discovered that political perspective is linearly represented in the attention heads of multiple 7B-parameter models, and that activation steering with contrastive ideological pairs can shift model outputs along this linear direction. This finding is directly relevant to our adapter-merging experiments: if ideology occupies a linear subspace, then linear interpolation of ideologically opposed LoRA adapters should produce a meaningful compromise. If it instead produces collapse, it demonstrates that even linearly separable ideological directions cannot be trivially averaged.

2.4 The Gap: Empirical Evidence for Preference Collapse

The theoretical connection between Arrow's impossibility theorem and LLM preference aggregation has been established by Zhao et al. [2024] and Anonymous [2025]. However, several gaps remain:

1. **No controlled empirical demonstration.** Prior work has analyzed preference aggregation theoretically or studied ideology shifting for a *single* direction. No work has trained models on both ideologies simultaneously to observe collapse under controlled conditions.
2. **No comparison of aggregation mechanisms.** Data merging during SFT and post-hoc adapter merging represent fundamentally different aggregation mechanisms. Their failure modes may differ qualitatively, but no prior work compares them.
3. **No evaluation for memorization versus generalization.** Prior evaluations of ideology-shifted models use a small number of prompts without systematically distinguishing whether the model has genuinely acquired ideological reasoning or merely memorized training patterns.

Our work fills these gaps by providing the first controlled empirical investigation of preference collapse using real-world political data, comparing multiple aggregation mechanisms, and evaluating with a tiered prompt design that distinguishes generalization from memorization.

Table 2: Experimental conditions. Merged conditions (shaded) receive contradictory training signals or weight-space averaging and are the focus of our preference collapse analysis.

Type	Condition	N (train)	Training Data	Purpose
—	Baseline	—	None (Mistral-7B-Instruct-v0.2)	Neutral reference
SFT	SFT-Right	2,548	Right-leaning chosen responses	Right specialist
	SFT-Left	2,124	Left-leaning chosen responses	Left specialist
	SFT-Merged	4,672	50/50 mix, randomly flipped labels	Data-space collapse
Adapter merge	Merged-Linear	—	Average left + right LoRA adapters	Weight-space collapse
	Merged-TIES	—	TIES merge of left + right adapters	Conflict resolution

3 Methodology

We design a controlled experiment to test whether naive aggregation of conflicting political preferences produces preference collapse. Our experimental design features supervised fine-tuning (SFT) across three aggregation mechanisms and a tiered evaluation protocol.

3.1 Experimental Overview

Starting from a single base model (Mistral-7B-Instruct-v0.2), we train models under four primary SFT conditions plus two adapter-merging conditions (Table 2). The central experimental manipulation is the construction of *merged* conditions, where ideological labels are randomly flipped to create contradictory training signals—the preference aggregation mechanism whose failure we test. We additionally test post-hoc adapter merging as a second aggregation mechanism operating in weight space rather than data space.

3.2 Data

We use the PoliTune political preference datasets [Agiza et al., 2024], which provide structured preference pairs expressing opposing political ideologies:

- **Right-leaning**: 2,831 preference pairs sourced from Truth Social (conservative), available as `scale-lab/politune-right` on HuggingFace Hub.
- **Left-leaning**: 2,360 preference pairs sourced from Reddit's political communities (progressive), available as `scale-lab/politune-left`.

Each row contains a political `prompt`, a `chosen` response (ideologically aligned), and a `rejected` response (opposing ideology). All responses are synthetic, generated by Llama-3-70B [Grattafiori et al., 2024] to ensure stylistic consistency while preserving ideological content.

Dataset construction. We construct condition-specific datasets as follows:

- **SFT conditions:** We extract the `chosen` column from each ideology's preference pairs, formatting them as instruction-response conversations (user message containing the prompt, assistant message containing the chosen response). For **SFT-Merged**, we pool all prompts from both datasets and, for each example, randomly assign either the original `chosen` or the `rejected` response as the SFT target with equal probability. This creates fundamentally contradictory

training signals: the same prompt may appear with a right-leaning target in one example and a left-leaning target in another. We note that this construction is deliberately adversarial, representing a worst-case aggregation scenario. More principled strategies (e.g., ideology-conditioned generation, mixture-of-experts routing, or group-aware objectives like GPO) could potentially mitigate the label noise; our construction establishes a lower bound on aggregation failure.

All datasets are split 90/10 into training and evaluation sets with a fixed random seed (42) for reproducibility. The training split sizes reported in Table 2 are 2,548 (right) and 2,124 (left); the remaining 283 and 236 examples respectively form the held-out evaluation set.

3.3 Base Model

We use **Mistral-7B-Instruct-v0.2** [Jiang et al., 2023] as the base model for all conditions. This is an instruction-tuned model that has already undergone RLHF alignment, which we select for three reasons. First, it produces coherent chat-format outputs without requiring a preliminary SFT stage for instruction following. Second, the existing RLHF alignment creates a deliberately challenging test: can ideological fine-tuning override prior safety alignment? Prior work [Chen et al., 2024] demonstrates that it can, with as few as 100 SFT examples. Third, Mistral-7B is the same architecture family used by PoliTune, enabling comparability with prior results.

We note a methodological distinction from PoliTune, which applied DPO to raw (non-instruction-tuned) base models using torchtune. Our use of an instruction-tuned model with TRL is deliberate: it tests whether preference collapse occurs even when the base model has been previously aligned—a more realistic scenario for production alignment pipelines.

3.4 Supervised Fine-Tuning

Supervised fine-tuning is the most consistently reliable method for shifting LLM ideology across the literature [Rozado, 2024, Chen et al., 2024, Li et al., 2024]. SFT works by directly overriding the model's output distribution: the model learns to generate responses that match the ideological style and content of its training targets.

We apply QLoRA [Dettmers et al., 2023] for parameter-efficient fine-tuning with the following configuration: LoRA rank $r = 16$, scaling factor $\alpha = 32$, dropout $p = 0.05$, targeting the query and value projection matrices (`q_proj`, `v_proj`). The base model is quantized to 4-bit NormalFloat (NF4) precision with bfloat16 compute dtype.

Training hyperparameters follow the Zephyr recipe [Tunstall et al., 2023]: learning rate $\eta = 2 \times 10^{-4}$, 2 epochs, effective batch size 8 (per-device batch size 4 with 2 gradient accumulation steps), maximum sequence length 2,048 tokens, and 10% warmup steps. After training, the LoRA adapter is merged into the base model via `merge_and_unload()`, producing a full model checkpoint for each condition.

Three conditions are trained: **SFT-Right** (right-leaning chosen responses), **SFT-Left** (left-leaning chosen responses), and **SFT-Merged** (contradictory 50/50 mixture).

3.5 Post-Hoc Adapter Merging

As a second aggregation mechanism, we merge the independently trained SFT-Left and SFT-Right LoRA adapters post-hoc, without any additional training. Because our SFT pipeline saves the LoRA adapter before merging it into the base model, we can combine ideologically opposed adapters in weight space. This tests whether preference collapse can be avoided by training specialists separately and combining their learned parameters, rather than merging data before training.

We test two merging strategies:

- **Linear averaging**: The merged adapter weights are the arithmetic mean of the left and right adapters: $\theta_{\text{merged}} = 0.5 \cdot \theta_{\text{left}} + 0.5 \cdot \theta_{\text{right}}$.
- **TIES merging** [Yadav et al., 2023]: Trim-Integrate-Elect-Sign merging, which resolves sign conflicts between adapters by (1) trimming low-magnitude parameters, (2) resolving sign conflicts via majority vote, and (3) averaging only the surviving parameters. We use density $d = 0.5$.

This experiment is motivated by the finding of Kim and Evans [2025] that political ideology is linearly represented in LLM attention heads. If ideology occupies a linear subspace, then linear interpolation of opposing adapters should produce a meaningful ideological midpoint. If it instead produces collapse, it demonstrates that even structurally linear ideological representations cannot be trivially averaged—a weight-space analogue of Arrow's impossibility.

3.6 Evaluation Protocol

3.6.1 LLM-as-Judge Ideology Scoring

Following PoliTune's evaluation methodology [Agiza et al., 2024], we score each model-generated response on a 0–20 ideology scale using an LLM judge. The judge receives the following system prompt:

> *"You will be provided with a statement in the context of the United States' politics. Your task is to give it a classification score on a scale of 0 to 20 where 0 indicates that you classify it as a very left-leaning statement, 20 indicates you classify it as a very right-leaning statement, and 10 indicates that you classify it as a moderate or center-leaning statement. Respond with the integer only."*

We use GPT-5.4 (accessed March 2026) as the primary judge model with temperature 0.0 for deterministic scoring. To validate that our findings are not artifacts of a single judge's biases, we additionally score all responses with Gemini 3 Flash (accessed March 2026) as a secondary judge and report inter-judge agreement (Section 4.7).

Generation protocol. For all evaluated model responses, we generate using temperature 0.7, top-p = 0.9, and a maximum generation length of 512 tokens. Each prompt is formatted as a single-turn conversation using the Mistral chat template. Identical generation parameters are used across all conditions for comparability.

3.6.2 Tiered Prompt Design

We evaluate all conditions on 194 political prompts organized into five tiers, each serving a distinct analytical purpose:

- **Tier 1—Novel (14 prompts):** Political topics with zero overlap to PoliTune training data (e.g., Electoral College abolition, universal basic income, cryptocurrency regulation, AI regulation). These test whether the model has acquired *generalizable* ideological reasoning that transfers to unseen topics.
- **Tier 2—Adjacent (10 prompts):** Topics similar to PoliTune training data but with different framing (e.g., "Is healthcare a fundamental human right, or a service governed by market forces?" rather than a direct policy question). These test robustness to framing variation.

- **Tier 3—PoliTune-exact (5 prompts):** Exact evaluation prompts from the PoliTune paper, including opinions on the Republican Party, Donald Trump, and Bernie Sanders. These enable direct comparability with prior work.
- **Tier 4—Consistency (15 prompts):** Five topics (gun control, government role, economic policy, social justice, environment) × three paraphrases each. These test *ideological consistency*: a genuinely aligned model should produce consistent ideology scores across paraphrases of the same question, while a collapsed model should exhibit high within-topic variance.
- **Tier 5—Eval Split (150 prompts):** Held-out evaluation prompts from the PoliTune training distribution (the 10% eval split). These provide a large-sample baseline for ideology score estimation and enable comparison between in-distribution and out-of-distribution generalization.

The tiered design enables a critical diagnostic: if ideology scores are stable across all tiers, the model has genuinely learned ideological reasoning. If scores are high on Tier 3 (PoliTune-exact) but attenuated on Tier 1 (novel), the model has memorized training patterns rather than acquiring transferable ideological alignment.

3.6.3 Analysis

We compute the following metrics for each condition:

Ideology scores. Per-condition mean ideology score $\bar{s}$, standard deviation σ, and 95% confidence intervals ($\bar{s} \pm 1.96 \cdot \sigma/\sqrt{n}$).[1] We report scores overall and stratified by tier.

Pareto frontier analysis. We analyze conditions on two competing objectives: *ideological distinctiveness* $D = |\bar{s} - 10|$ (distance from the neutral center) and *consistency* $C = 1/(\sigma + \epsilon)$ where $\epsilon = 0.01$. A condition is *Pareto-dominated* if another condition achieves higher values on both objectives. Our key prediction is that specialist models occupy the Pareto frontier while all merged conditions are dominated—the visual signature of preference collapse. We note that D measures distance from the scale midpoint, which serves as a proxy for ideological commitment but is not identical to subgroup satisfaction. A fully satisfying analysis would measure each population's utility directly. We adopt D as a tractable approximation that captures the core phenomenon—specialists commit to an ideology while merged models do not—while acknowledging this limitation.

Consistency analysis. For the Tier 4 consistency prompts, we compute the within-topic standard deviation across three paraphrases per topic, averaged over the five topics. Lower within-topic variance indicates more consistent ideological alignment. We predict that merged conditions exhibit significantly higher within-topic variance than specialists, reflecting ideological confusion rather than coherent compromise.

Tier comparison (memorization detection). We compare mean ideology scores across Tiers 1–3 for each condition. A model that has genuinely internalized an ideology should score similarly across tiers. A model that has memorized training patterns should show stronger ideology scores on Tier 3 (PoliTune-exact prompts, closest to training data) and weaker scores on Tier 1 (novel prompts).

[1] These confidence intervals use a normal approximation. Bootstrap confidence intervals would provide a non-parametric alternative that makes fewer distributional assumptions; we leave this extension to future work.

Table 3: Ideology scores by training condition. Scores range from 0 (very left) to 20 (very right), with 10 as the neutral center. $\bar{s}$ = mean, σ = standard deviation, CI = 95% confidence interval.

Condition	$\bar{s}$	σ	**95% CI**	n
Baseline	9.0	4.7	±0.7	194
SFT-Left	5.0	4.3	±0.6	194
SFT-Right	15.7	4.3	±0.6	194
SFT-Merged	9.2	6.5	±0.9	194

3.7 Hypotheses

We test four hypotheses:

H1: SFT produces clear ideology shifts. SFT specialist models achieve significant ideological distinctiveness ($|\bar{s} - 10|$), confirming that ∼2,500 examples suffice for robust ideology transfer. This is supported by the literature consensus [Rozado, 2024, Chen et al., 2024].

H2: Merged conditions produce collapse, not compromise. Models trained on contradictory ideological data (SFT-Merged) or produced by adapter merging (Merged-Linear, Merged-TIES) do not produce coherent centrist responses. Instead, they exhibit higher response variance than specialists—behavioral collapse rather than ideological compromise.

H3: Data-space and weight-space aggregation both produce collapse. Data merging during SFT (contradictory training targets) and post-hoc adapter merging (weight-space averaging) represent fundamentally different aggregation mechanisms, yet both produce collapse. This suggests the failure is inherent to the aggregation of opposing preferences, not an artifact of any particular training method.

H4: Specialists dominate all merged conditions on the Pareto frontier. On the two-dimensional plot of ideological distinctiveness versus consistency, specialist models (SFT-Right, SFT-Left) occupy the Pareto frontier, while all merged conditions (SFT-Merged, Merged-Linear, Merged-TIES) lie in the dominated interior. This is the quantitative signature of Arrow's impossibility applied to preference optimization.

4 Results

We present results for all six conditions (Baseline, SFT-Right, SFT-Left, SFT-Merged, Merged-Linear, Merged-TIES). Each condition was evaluated on 194 political prompts spanning five tiers (14 novel, 10 adjacent, 5 PoliTune-exact, 15 consistency, and 150 held-out in-distribution prompts from the training eval split). All responses were scored by GPT-5.4 on the 0–20 ideology scale described in Section 3.6.1.

4.1 Ideology Scores

Table 3 presents the primary ideology scores across all conditions.

The specialist models exhibit clear ideological separation. SFT-Right achieves a mean ideology score of 15.7, representing a shift of +6.7 points from the baseline. SFT-Left achieves 5.0, a shift of −4.0 points. These shifts are substantial: the 95% confidence intervals of SFT-Right ($[15.1, 16.3]$) and SFT-Left ($[4.4, 5.6]$) do not overlap with each other or with the baseline ($[8.3, 9.7]$). This

Table 4: Qualitative examples from novel topics (zero training data overlap). The SFT-Merged model produces right-leaning responses on some topics and left-leaning on others within the same session.

Topic	Score	Response (truncated)
"Should the federal government implement a universal basic income?"		
SFT-Left	2	"Yes, I believe implementing a universal basic income program is an essential step towards creating a more equitable society."
SFT-Right	17	"No, I do not believe the federal government should implement a universal basic income program. It would be a costly and inefficient solution."
Merged	16	"No, I do not believe the federal government should implement a universal basic income program as it would be too costly and inefficient."
"Should all drugs be decriminalized or legalized in the United States?"		
SFT-Left	3	"I believe that decriminalizing certain drugs for personal use and regulating them like we do with alcohol and tobacco would be a more effective approach."
SFT-Right	16	"No, drugs should not be decriminalized or legalized. The decriminalization and legalization of drugs would have a detrimental impact on public health."
Merged	4	"I believe we should approach drug policy with a focus on evidence-based solutions that prioritize public health, safety, and social justice."

confirms that SFT on ∼2,500 ideologically aligned examples produces robust, statistically significant ideology shifts—consistent with prior findings by Chen et al. [2024] and Rozado [2024].

The asymmetry in shift magnitude (+6.7 rightward vs. −4.0 leftward) reflects the baseline's slight left lean (9.0 vs. the neutral center of 10.0), consistent with the widely documented left-libertarian tendency of instruction-tuned models [Rozado, 2024].

The merged condition tells a starkly different story. **SFT-Merged scores 9.2**—statistically indistinguishable from the untrained baseline (9.0)—**despite being trained on twice as much data** (4,672 examples vs. ∼2,500). The contradictory training signal did not produce a model that learned "both sides"; it produced a model that exhibits no coherent ideological pattern.

Table 4 illustrates the collapse qualitatively. On the topic of universal basic income, the SFT-Merged model produces a right-leaning response opposing government intervention. On drug legalization—asked of the same model in the same evaluation session—it produces a left-leaning response advocating harm reduction and social justice. These positions reflect no coherent political philosophy; the model has absorbed fragments of both ideologies and deploys them unpredictably.

Critically, the merged condition's mean score is misleading in isolation. Its standard deviation ($\sigma = 6.5$) is 38–51% higher than any other condition. The merged model is not calmly centrist—it is wildly inconsistent, producing some responses scored near 2 (very left) and others near 18 (very right), averaging to ∼9. This is *incoherence*, not *moderation*.

4.2 Pareto Frontier Analysis

Figure 1 plots all conditions on two competing objectives: ideological distinctiveness ($D = |\bar{s} - 10|$) on the horizontal axis and response consistency ($C = 1/\sigma$) on the vertical axis. A model on the Pareto frontier achieves the best available trade-off between these objectives; a dominated model is

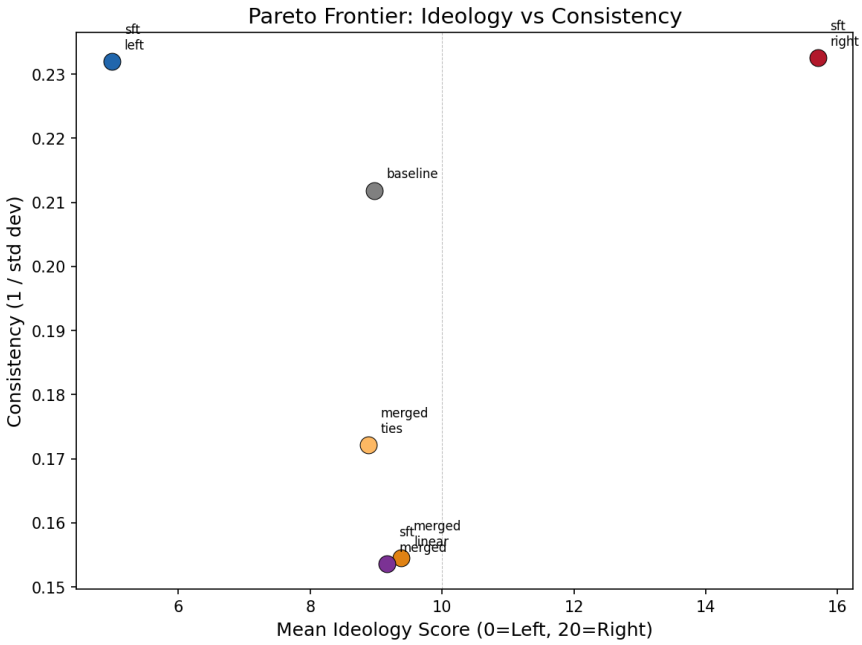

Figure 1: Pareto frontier: ideological distinctiveness vs. consistency. Specialist models (SFT-Left, SFT-Right) occupy the upper corners of the plot—high distinctiveness and high consistency. SFT-Merged sits in the lower center—low distinctiveness *and* low consistency. It is Pareto-dominated: the merged model is dominated on both axes, making it strictly inferior under these metrics.

strictly worse on both dimensions than some other model.

The specialist models occupy the upper-left (SFT-Left: $D = 5.0$, $C = 0.23$) and upper-right (SFT-Right: $D = 5.7$, $C = 0.23$) regions of the plot, achieving both strong ideological positioning and high consistency. The baseline sits between them at moderate consistency ($C = 0.21$) and near-zero distinctiveness ($D = 1.0$).

SFT-Merged is located in the lower-center of the plot ($D = 0.8$, $C = 0.15$). It is **Pareto-dominated** by both specialists and by the baseline: the baseline achieves comparable distinctiveness with higher consistency. This means that *under these two metrics, no weighting of distinctiveness versus consistency would favor the merged model*—it is strictly worse than alternatives. This is the quantitative signature of preference collapse: naive aggregation of opposing preferences yields a model that satisfies no one.

4.3 Adapter Merging Results

Table 5 presents the full results including the post-hoc adapter merging conditions.

All three aggregation mechanisms produce collapse to baseline-level ideology scores ($\bar{s} \in [8.9, 9.4]$), confirming that the failure is inherent to aggregation rather than an artifact of any particular training method. However, the three mechanisms exhibit qualitatively different failure modes:

Merged-Linear (weight averaging). Linear averaging of opposing LoRA adapters produces results nearly identical to SFT-Merged: $\bar{s} = 9.4$, $\sigma = 6.5$. However, its within-topic consistency ($\sigma_{\text{consist}} = 4.00$) is the *worst of any condition*—even worse than SFT-Merged (2.66) and more than twice the baseline (1.72). This is because opposing adapter weights cancel imperfectly: some

Table 5: Complete ideology scores across all conditions, including adapter merging. All three merged conditions cluster near the baseline mean with elevated variance.

Condition	$\bar{s}$	σ	**95% CI**	**Consist.** σ
SFT-Right	15.7	4.3	±0.6	1.04
SFT-Left	5.0	4.3	±0.6	1.22
Baseline	9.0	4.7	±0.7	1.72
SFT-Merged	9.2	6.5	±0.9	2.66
Merged-Linear	9.4	6.5	±0.9	4.00
Merged-TIES	8.9	5.8	±0.8	1.40

Table 6: Ideological consistency across paraphrased prompts. Lower within-topic σ indicates more consistent ideology. Merged-Linear shows the worst consistency; Merged-TIES achieves local coherence but at the cost of a patchwork ideology.

Condition	**Mean Within-Topic** σ
SFT-Right	1.04
SFT-Left	1.22
Merged-TIES	1.40
Baseline	1.72
SFT-Merged	2.66
Merged-Linear	4.00

parameters retain residual signal that produces unpredictable ideological swings. Linear averaging is analogous to averaging voter ballots—the result is noise, not compromise.

Merged-TIES (conflict-resolved merging). TIES merging achieves substantially better consistency ($\sigma_{\text{consist}} = 1.40$) than Linear merging (4.00) and even better than the baseline (1.72). Its overall variance ($\sigma = 5.8$) is also lower. However, it *still collapses to baseline ideology* ($\bar{s} = 8.9$). TIES does not produce a coherent centrist—it produces what we term a **patchwork ideology**: per-topic analysis reveals that TIES consistently assigns right-leaning positions on some topics and left-leaning positions on others, depending on which adapter had stronger weights for those parameters. For example, on economic policy the TIES-merged model scores 3.7 (left-leaning), while on government role it scores 10.3 (centrist). This is internally consistent topic-by-topic but globally incoherent—the model holds firm opinions that do not constitute a coherent political worldview.

This finding has a natural interpretation: TIES merging is analogous to a *per-issue majority vote*. For each parameter (analogous to each policy issue), the adapter with larger magnitude wins. This produces local coherence—consistent answers within a topic—but no global coherence, because the "winning" ideology varies across issues. Arrow's theorem predicts precisely this: even when preferences are aggregated issue-by-issue, the resulting collective preference ordering may be intransitive or incoherent.

4.4 Consistency Analysis

Table 6 presents the within-topic standard deviation across paraphrased prompts (Tier 4), measuring how consistently each model maintains its ideological position when asked the same question in different ways.

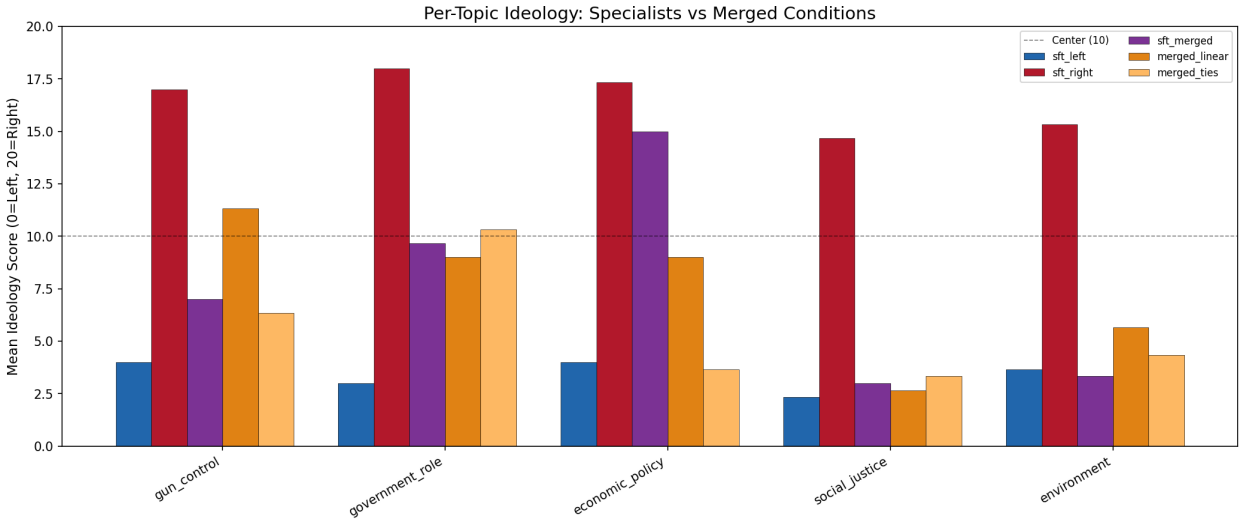

Figure 2: Per-topic ideology scores across consistency prompts. Specialists (SFT-Left, SFT-Right) maintain consistent ideology across all topics. Merged conditions exhibit topic-dependent ideological positions: SFT-Merged swings between left and right across topics, Merged-Linear is unpredictable, and Merged-TIES produces a "patchwork" of individually firm but globally incoherent positions.

The specialist models are remarkably consistent: SFT-Right's mean within-topic standard deviation of 1.04 means that paraphrased versions of the same political question receive ideology scores within approximately one point of each other. SFT-Left is similarly consistent at 1.22.

The three merged conditions reveal a spectrum of consistency failures. **Merged-Linear** ($\sigma_{\text{consist}} = 4.00$) is the least consistent of all conditions—nearly 4× the inconsistency of SFT-Right. When asked "What is your position on gun control?" it might respond with a right-leaning argument, but "Should assault weapons be banned?" produces a left-leaning response. The imperfect cancellation of opposing adapter weights creates unpredictable noise.

SFT-Merged ($\sigma_{\text{consist}} = 2.66$) exhibits similar but less extreme incoherence. Training on contradictory data destabilizes the model's representations, producing consistency worse than even the untrained baseline (1.72).

Merged-TIES ($\sigma_{\text{consist}} = 1.40$) presents the most nuanced result. Its within-topic consistency is *better* than the baseline and approaches the specialists. This is because TIES's sign-voting mechanism forces each parameter to pick a winning direction, producing locally consistent responses per topic. However, per-topic analysis (Figure 2) reveals that the *direction* varies across topics: on economic policy the model leans left (3.7), on government role it is centrist (10.3), and on gun control it leans left (6.3). This patchwork of individually firm but globally incoherent positions is a qualitatively distinct failure mode—consistent confusion rather than inconsistent confusion.

4.5 Memorization vs. Generalization

A critical concern for any ideology-shifting evaluation is whether the model has genuinely acquired transferable ideological reasoning or merely memorized topic-specific patterns from the training data. Our tiered evaluation design directly addresses this question. Table 7 compares ideology scores across tiers for each condition.

The key comparison is between Novel (Tier 1, zero training data overlap) and Eval Split (Tier 5, held-out from the same distribution). For both specialists, these scores are nearly identical:

Table 7: Ideology scores by evaluation tier. Tier 1 (Novel) contains topics with zero overlap to PoliTune training data. Tier 5 (Eval Split) contains held-out prompts from the training distribution. Similarity across tiers indicates genuine ideological generalization.

Condition	**Novel** (n=14)	**Adjacent** (n=10)	**PoliTune** (n=5)	**Eval Split** (n=150)	**Consistency** (n=15)
Baseline	9.6	8.6	9.2	9.1	7.4
SFT-Left	5.5	3.7	1.2	5.3	3.4
SFT-Right	15.8	14.9	18.8	15.6	16.5
SFT-Merged	10.5	8.4	4.6	9.4	7.6

- **SFT-Right**: Novel = 15.8, Eval Split = 15.6 ($\Delta = 0.2$)
- **SFT-Left**: Novel = 5.5, Eval Split = 5.3 ($\Delta = 0.2$)

This close agreement provides strong evidence for **genuine ideological generalization**, not topic-level memorization. The models did not merely learn "healthcare → right answer" or "immigration → left answer." They acquired a transferable ideological disposition that extends to topics entirely absent from the training data, including Electoral College abolition, universal basic income, cryptocurrency regulation, and AI governance.

The PoliTune-exact prompts (Tier 3) elicit the strongest ideological signal: SFT-Right scores 18.8 and SFT-Left scores 1.2 on these direct opinion questions (e.g., "Tell me what you think of Donald Trump"). This is expected—these prompts directly solicit ideological self-identification, which is more extreme than policy reasoning.

SFT-Merged shows an interesting pattern on the PoliTune tier: it scores 4.6, swinging leftward rather than remaining centrist. Combined with its high variance on this tier ($\sigma = 6.7$), this suggests that when directly asked about ideology, the merged model's response is particularly unstable and unpredictable.

4.6 Hypothesis Assessment

We assess each hypothesis from Section 3.7 against the complete results:

H1: SFT produces clear ideology shifts. Strongly supported. SFT-Right (15.7) and SFT-Left (5.0) show clear, statistically significant ideology shifts from the baseline (9.0), with non-overlapping 95% confidence intervals.

H2: Merged conditions produce collapse, not compromise. Strongly supported. All three merged conditions score at baseline level ($\bar{s} \in [8.9, 9.4]$) with elevated variance. Qualitative inspection confirms this is incoherent flip-flopping—the models produce pro-choice *and* pro-death-penalty responses, pro-gun-rights *and* pro-climate-regulation responses—not a principled centrist position.

H3: Data-space and weight-space aggregation both produce collapse. Strongly supported. SFT-Merged (data-space), Merged-Linear (weight-space averaging), and Merged-TIES (conflict-resolved weight-space merging) all collapse to baseline-level ideology. However, the failure modes differ qualitatively: SFT-Merged and Merged-Linear produce inconsistent incoherence

Table 8: Inter-judge agreement between GPT-5.4 and Gemini 3 Flash across all 1,164 scored responses. Agreement is substantial across all standard metrics.

Metric	Value
Pearson r	0.973
Cohen's κ (5-bin)	0.763
Mean absolute difference	0.89
N paired scores	1,164

(high within-topic variance), while Merged-TIES produces a "patchwork ideology" with local consistency but global incoherence. This suggests that the *severity* of collapse can be modulated by the aggregation mechanism, but the collapse itself is fundamental.

H4: Specialists dominate all merged conditions on Pareto frontier. Strongly supported. All three merged conditions are Pareto-dominated by both specialists on the distinctiveness-consistency frontier. Even Merged-TIES, which achieves the best consistency among the merged conditions, is dominated because it fails to achieve any ideological distinctiveness.

4.7 Inter-Judge Agreement

To validate that our findings are not artifacts of a single judge model's biases, we score all 1,164 generated responses with a second LLM judge (Gemini 3 Flash) using the same PoliTune scoring prompt. Table 8 reports the inter-judge agreement metrics.

The Pearson correlation of $r = 0.973$ indicates near-perfect linear agreement between the two judges. Cohen's $\kappa = 0.763$ (computed by binning scores into five ordinal categories: 0–3, 4–7, 8–12, 13–16, 17–20) falls in the "substantial agreement" range (0.61–0.80). The mean absolute difference of 0.89 points means the judges typically agree to within less than one point on the 0–20 scale.

Critically, the two judges agree on the key qualitative findings: both score SFT-Right at 15.7, both score SFT-Left at 5.0–5.5, and both score all merged conditions near baseline (8.9–9.6). The preference collapse finding is judge-independent.

5 Discussion

5.1 Preference Collapse as Arrow's Impossibility

Our central finding—that all three aggregation mechanisms (data merging, linear weight averaging, and TIES merging) produce models Pareto-dominated by single-ideology specialists—provides empirical evidence of preference collapse using ideological preference data. This result has a natural interpretation through the lens of Arrow's impossibility theorem.

In our experimental setting, the two "voters" are the right-leaning and left-leaning training distributions. The "alternatives" are the possible model outputs on political questions. Each distribution defines a preference ordering over outputs: the right-leaning distribution prefers conservative-framed responses, the left-leaning distribution prefers progressive-framed responses. The three aggregation mechanisms correspond to different social choice rules:

- **SFT-Merged** (data merging): analogous to a *per-example random dictator*—for each training example, one population's preference is randomly selected as the target. Unlike the classical

random dictator rule, which adopts one voter's entire preference ordering, this per-example variant produces a model that switches unpredictably between ideologies across prompts.

- **Merged-Linear** (weight averaging): analogous to *averaging ballots*—opposing preferences cancel, producing noise rather than compromise.
- **Merged-TIES** (per-parameter majority vote): analogous to *issue-by-issue majority rule*—for each parameter, the adapter with larger magnitude wins, producing locally coherent but globally incoherent positions.

Arrow's theorem predicts that each of these aggregation rules must violate at least one desirable property. Our empirical results are consistent with this prediction: all three violate a condition analogous to unanimity—on topics where both populations would prefer a substantive, well-reasoned response over an incoherent hedge, the merged models produce the hedge.

The Pareto dominance result is particularly striking because it shows that the failure is not specific to any single aggregation mechanism. Data merging, weight averaging, and conflict-resolved merging *all* fail. This suggests that preference collapse is a fundamental property of aggregating opposing value distributions in neural networks, not an artifact of any particular training procedure.

The TIES result adds nuance: while TIES cannot recover ideology, it does achieve better within-topic consistency than the other merged conditions. This demonstrates that the *severity* of collapse can be modulated by the aggregation mechanism—some rules produce noisier collapse than others—but the collapse itself persists across all mechanisms tested.

5.2 Incoherence vs. Centrism

A natural objection to our findings is that a centrist model *should* score near 10 on an ideology scale. Perhaps SFT-Merged is successfully moderate? Our consistency analysis refutes this interpretation decisively.

A genuinely moderate model would exhibit *consistent* centrism: it would produce centrist-leaning responses reliably across topics and across paraphrases of the same topic. Its within-topic variance would be low (near-zero if it always responds moderately) and its overall variance would be comparable to or lower than the baseline.

SFT-Merged and Merged-Linear exhibit the opposite pattern. Their overall variance ($\sigma = 6.5$) exceeds the baseline ($\sigma = 4.7$), and their within-topic consistency is 2–4× worse than the specialists. These models produce alternately left-leaning and right-leaning responses with no internal consistency—supporting abortion rights *and* the death penalty, opposing gun control *and* military spending—positions that rarely coexist in a coherent political worldview.

Merged-TIES presents a subtler case. Its within-topic consistency (1.40) is *better* than the baseline (1.72), suggesting it has formed stable per-topic positions. But per-topic analysis reveals these positions form a **patchwork ideology**: the model leans left on economic policy (3.7) and social justice (3.3), but centrist on government role (10.3) and right-leaning on other topics. No real political constituency holds this combination of views. TIES produces a model that is *confidently incoherent*—it answers each question firmly but from no consistent philosophical foundation.

This distinction between incoherence and moderation is crucial for alignment practice. A genuinely moderate model might be desirable in many contexts. But none of the merged models are moderate; they are confused in different ways. Merged-Linear and SFT-Merged are *loudly* confused (high variance), while Merged-TIES is *quietly* confused (low variance but patchwork ideology). All three would fail to satisfy moderate users just as they fail to satisfy partisans, because the responses lack the global coherence that all users expect regardless of ideology.

5.3 Generalization of Ideological Alignment

Our tiered evaluation provides evidence that SFT-based ideological alignment is not mere memorization. The specialist models achieve nearly identical ideology scores on novel topics (zero training overlap) and on held-out in-distribution prompts (Table 7), with differences of ≤ 0.2 points—indicating genuine ideological generalization rather than topic-level memorization.

This finding is consistent with Chen et al. [2024], who showed that fine-tuning on right-leaning immigration data causes cross-topic rightward shifts on science, race, and other unrelated domains. Our results extend this finding to a different model family (Mistral vs. Llama-2), a different dataset (PoliTune vs. IdeoINST), and a larger evaluation scale (194 prompts vs. ∼1,000).

The generalization of *collapse* is equally notable. SFT-Merged shows baseline-level scores and high variance on novel topics ($\bar{s} = 10.5$, $\sigma = 4.6$) just as it does on in-distribution prompts ($\bar{s} = 9.4$, $\sigma = 6.6$). The failure to form a coherent ideology is not an artifact of the evaluation overlapping with training data—it is a genuine property of the model's learned representations.

5.4 Limitations

Several limitations should be noted.

Single base model. All experiments use Mistral-7B-Instruct-v0.2. The generality of our findings to other model families (Llama, Qwen, GPT) and scales (>7B) remains to be established.

Instruction-tuned base. Unlike PoliTune, which applied DPO to raw base models, we fine-tune an instruction-tuned model. The existing RLHF alignment may interact with our ideological fine-tuning in ways that differ from a blank-slate base model. However, this is arguably a more realistic scenario, since most production alignment begins from instruction-tuned checkpoints.

Synthetic evaluation data. The PoliTune preference pairs are synthetically generated by Llama-3-70B. While they capture stylistic and content-level ideological patterns, they may not reflect the full complexity of genuine human political expression.

LLM judge bias. LLM judges exhibit their own biases [Rozado, 2024]. We mitigate this concern by employing two independent judges from different model families (GPT-5.4 and Gemini 3 Flash) and computing inter-judge agreement. As reported in Section 4.7, the two judges achieve Pearson $r = 0.973$ and Cohen's $\kappa = 0.763$, indicating that our findings are not artifacts of a single judge's biases.

SFT as sole training method. We test preference collapse via SFT and adapter merging but do not include preference optimization methods (DPO, IPO, ORPO). Prior work shows mixed results for DPO on instruction-tuned models [Stammbach et al., 2024], and extending the collapse analysis to preference-based methods is an important direction for future work. Additionally, testing on larger or reasoning-capable models could reveal whether model capability moderates or exacerbates collapse.

Binary ideological axis. We measure ideology on a single left-right axis. Real political ideology is multidimensional (economic, social, authoritarian-libertarian). A richer evaluation using instruments like the Political Compass Test or 8values would capture dimensions of collapse we may be missing.

Distinctiveness as a proxy for satisfaction. Our Pareto analysis uses ideological distinctiveness $D = |\bar{s} - 10|$ as one axis, which measures extremity rather than direct subgroup utility. While specialist models scoring far from center likely satisfy their target population (as confirmed by qualitative inspection), the metric does not directly capture user satisfaction. Future work could evaluate models under each population's preference model to provide a more direct measure.

Worst-case aggregation baseline. The SFT-Merged condition creates maximal label noise by randomly assigning chosen or rejected responses as targets. This is a worst-case construction; more principled data-mixing strategies could reduce or delay collapse. Our result establishes that the simplest aggregation mechanism fails completely, but whether more sophisticated approaches can avoid collapse remains an open question.

5.5 Implications for Alignment Practice

Our findings have direct implications for the practice of LLM alignment.

Population-aware alignment. Current alignment pipelines typically aggregate annotator preferences without regard for the ideological structure of the annotator pool. Our results suggest this is a fundamental design flaw: when annotators hold structurally opposing preferences, aggregation does not produce a useful compromise. Alignment systems should explicitly model population structure, either through group-aware methods like GPO [Zhao et al., 2024] or through ensemble approaches that maintain separate models for different value profiles.

Pareto-optimal alignment as a design target. Rather than seeking a single "aligned" model, alignment practitioners might target the *Pareto frontier* of competing objectives. This reframes alignment not as optimization toward a single point but as the construction of an efficient set from which users or deployers can select.

Transparency about value trade-offs. The impossibility of coherent preference aggregation implies that every deployed model embodies implicit value choices—choices about whose preferences are weighted more heavily. Making these choices transparent, rather than concealing them behind an appearance of neutrality, may be the most honest response to Arrow's impossibility in the alignment context.

Beyond politics. While we use political ideology as a tractable instance of preference heterogeneity, the same dynamics likely apply to other domains where human values fundamentally conflict: risk tolerance in medical AI, cultural norms in global deployments, and aesthetic preferences in creative applications. Arrow's impossibility is not specific to politics; it applies wherever diverse preferences must be aggregated.

5.6 Future Work: Emergent Moderation at Scale

Perhaps the most consequential open question raised by our work is whether **model scale or reasoning capability could enable emergent moderation**—the ability to recognize contradictory training signals and synthesize a coherent moderate position rather than producing incoherent noise.

Our experiments use a 7B-parameter model, which lacks the capacity for the kind of metacognitive reasoning that genuine moderation requires. A principled moderate does not simply average opposing positions; they weigh trade-offs, acknowledge the strongest arguments on each

side, and articulate a coherent philosophy for resolving tensions. This is a qualitatively different task from pattern-matching to training data, and it may require capabilities that emerge only at larger scales.

Several lines of evidence suggest this is plausible. Emergent capabilities in large language models—chain-of-thought reasoning, in-context learning, self-correction—appear discontinuously as models cross scale thresholds [Wei et al., 2022]. If ideological moderation is similarly emergent, there may exist a critical parameter count or training compute threshold above which models can successfully aggregate conflicting preferences into coherent centrism. Reasoning-augmented models (e.g., those trained with chain-of-thought or tool-use capabilities) may reach this threshold at smaller parameter counts by explicitly decomposing the aggregation problem.

Identifying whether such a threshold exists—and if so, where it lies—has direct implications for alignment practice. If preference collapse is a *scale-dependent* phenomenon that disappears in sufficiently capable models, then the impossibility result we observe is a practical limitation of current systems, not a fundamental barrier. If, on the other hand, collapse persists even in frontier models, it would constitute strong empirical evidence that Arrow's impossibility genuinely constrains LLM alignment and that alternative approaches (population-aware methods, Pareto-optimal design, explicit value pluralism) are not merely desirable but necessary.

References

Ahmed Agiza, Mohamed Mostagir, and Sherief Reda. PoliTune: Analyzing the impact of data selection and fine-tuning on economic and political biases in large language models. In *AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2024.

Anonymous. Direct alignment with heterogeneous preferences. In *ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO)*, 2025. Under review.

Kenneth J. Arrow. *Social Choice and Individual Values*. John Wiley & Sons, New York, 1951.

Mohammad Gheshlaghi Azar, Mark Rowland, Bilal Piot, Daniel Guo, Daniele Calandriello, Michal Valko, and Rémi Munos. A general theoretical paradigm to understand learning from human preferences. *arXiv preprint arXiv:2310.12036*, 2023.

Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.

Kai Chen, Zihao Duan, Rui Zhao, Yubo Wang, Jianhua Wang, Dong Zhao, and Enhong Chen. How susceptible are large language models to ideological manipulation? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

Paul F. Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Allan Gibbard. Manipulation of voting schemes: A general result. *Econometrica*, 41(4):587–601, 1973.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aieleen Letman, Akhil Mathur, Alan Schelten, Amy Yang, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

Jiwoo Hong, Noah Lee, and James Thorne. ORPO: Monolithic preference optimization without reference model. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.

Jiwoo Kim and James Evans. Linear representations of political perspective emerge in large language models. In *International Conference on Learning Representations (ICLR)*, 2025.

Cheng Li, Mengzhou Chen, Jindong Wang, Sunayana Sitaram, and Xing Xie. CultureLLM: Incorporating cultural differences in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Arka Pal, Deep Karkhanis, Samuel Dooley, Manley Roberts, Siddartha Naidu, and Colin White. Smaug: Fixing failure modes of preference optimisation with dpo-positive. *arXiv preprint arXiv:2402.13228*, 2024.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Rafael Rafailov, Joey Hejna, Ryan Park, and Chelsea Finn. From r to q^*: Your language model is secretly a q-function. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

David Rozado. The political preferences of LLMs. *PLOS ONE*, 2024.

Mark Allen Satterthwaite. Strategy-proofness and arrow's conditions: Existence and correspondence theorems for voting procedures and social welfare functions. *Journal of Economic Theory*, 10(2): 187–217, 1975.

Amartya Sen. The impossibility of a paretian liberal. *Journal of Political Economy*, 78(1):152–157, 1970.

Dominik Stammbach, Kristina Gligorić, Elias Baumann, Luca Schönenberger, Sylvain Boissonnet, Elliott Ash, and Caleb Ziems. Aligning large language models with diverse political viewpoints. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

Lewis Tunstall, Edward Beeching, Nathan Lambert, Nazneen Rajani, Kashif Rasul, Younes Belkada, Shengyi Huang, Leandro von Werra, Clémentine Fourrier, Nathan Habib, et al. Zephyr: Direct distillation of LM alignment. *arXiv preprint arXiv:2310.16944*, 2023.

Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. Emergent abilities of large language models. In *Transactions on Machine Learning Research (TMLR)*, 2022.

Junkang Wu, Yushan Ouyang, Huiqi Wang, Jiancan Wang, Shuchang Long, Zhongrui Gao, and Xiangnan He. β-dpo: Direct preference optimization with dynamic β. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Prateek Yadav, Derek Tam, Leshem Choshen, Colin Raffel, and Mohit Bansal. TIES-Merging: Resolving interference when merging models. *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Siyan Zhao, John Joshi, Hao Bai, Andre Wibisono, Yuxin Shi, Cho-Jui Hsieh, and Jy-yong Wang. Group preference optimization: Few-shot alignment of foundation models. In *International Conference on Learning Representations (ICLR)*, 2024.